

Contents

1 Headhunter AI - Comprehensive Project Presentation	1
1.1 Enterprise Applicant Tracking System with AI Intelligence	1
1.2 Table of Contents	1
1.3 Executive Summary	2
1.4 System Architecture	2
1.5 Feature Deep Dive	4
1.6 Security Architecture	12
1.7 CI/CD & Testing Strategy	13
1.8 Roadmap & Future Features	16
1.9 Technical Onboarding	18
1.10 Business Value Proposition	21
1.11 Additional Resources	23
1.12 Key Takeaways	23
1.13 Next Steps	24
1.14 Conversion Instructions	25

1 Headhunter AI - Comprehensive Project Presentation

1.1 Enterprise Applicant Tracking System with AI Intelligence

Version: 1.11.0

Date: December 2025

Prepared By: Engineering Team

1.2 Table of Contents

1. [Executive Summary](#)
 2. [System Architecture](#)
 3. Feature Deep Dive
 4. [Security & Authentication](#)
 5. CI/CD & Testing Strategy
 6. Roadmap & Future
 7. Technical Onboarding
 8. Business Value
-

1.3 Executive Summary

1.3.1 What is Headhunter AI?

Headhunter AI is a production-grade, self-hosted Applicant Tracking System engineered for high-performance recruitment teams.

1.3.2 Key Differentiators

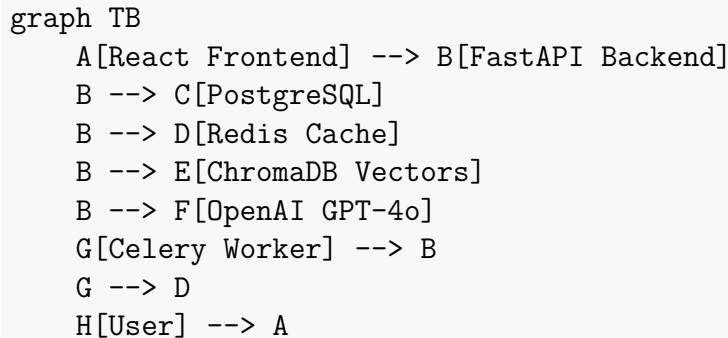
Traditional ATS	Headhunter AI
Manual data entry	AI-powered parsing
Basic file storage	Semantic search & matching
Generic job posts	Context-aware descriptions
Single-company	Multi-tenant architecture
Cloud subscription	Self-hosted, full control

1.3.3 Business Value

- **80% reduction** in manual data entry
- **60% faster** candidate screening
- **Complete data ownership** - no vendor lock-in
- **Unlimited users** - no per-seat pricing
- **Enterprise security** - self-hosted on your infrastructure

1.4 System Architecture

1.4.1 High-Level Overview



1.4.2 Technology Stack

Layer	Technology	Version	Purpose
Frontend	React	18	UI/UX

Layer	Technology	Version	Purpose
Styling	Tailwind CSS	3.x	Design System
Build Tool	Vite	5.x	Fast development
Backend	FastAPI	0.111+	REST API
Runtime	Python	3.13	Application logic
Database	PostgreSQL	15	Relational data
Cache	Redis	7	Caching & queues
Vector DB	ChromaDB	0.4+	Semantic search
AI Engine	OpenAI	GPT-4o-mini	Parsing & matching
Worker	Celery	5.3+	Background jobs
Deployment	Docker	Latest	Containerization

1.4.3 Project Structure

```

Headhunter/
  backend/          # FastAPI Application
    app/
      api/v1/        # REST Endpoints
        auth.py      # Authentication & SSO
        cv.py        # CV upload & parsing
        jobs.py      # Job management
        profiles.py  # Candidate profiles
        interviews.py # Interview scheduling
        analytics.py # Dashboard metrics
        sso.py        # SSO providers
      core/          # Config & utilities
        database.py  # DB connection
        security.py  # JWT & hashing
        email.py     # Email service
        cache.py     # Redis caching
      models/        # SQLAlchemy models
      services/      # Business logic
        ai_parser.py # CV parsing
        chroma.py    # Vector search
        matcher.py   # Candidate matching
    alembic/        # Database migrations
    tests/          # Unit & E2E tests
    scripts/        # Utility scripts
  frontend/        # React Application
    src/
      components/  # Reusable UI components
      pages/       # Page-level components
      hooks/       # Custom React hooks
      services/    # API client

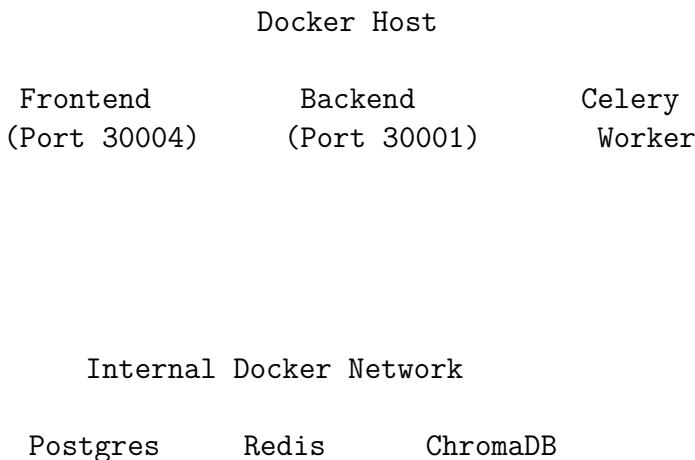
```

```

cypress/          # E2E tests
data/             # Persistent volumes
  db/            # PostgreSQL data
  raw/           # Uploaded CVs
docs/             # Documentation
.github/workflows/ # CI/CD pipelines

```

1.4.4 Deployment Architecture



Port Mapping:

- 30001 → Backend API
- 30002 → PostgreSQL (dev access)
- 30003 → ChromaDB (AI search)
- 30004 → Frontend UI
- 6380 → Redis (dev access)

1.5 Feature Deep Dive

1.5.1 1. Multi-Company Support

Business Value: One deployment serves unlimited companies with complete data isolation.

Features:

- Automatic company creation from email domain
- First user from domain becomes Admin
- Strict logical isolation (ALL queries filtered by `company_id`)
- Super Admin dashboard for cross-company management

Technical Implementation:

- Database: company_id foreign key on all major tables
- Security: Middleware ensures users only access their company data
- Performance: Indexed company_id columns for fast filtering

Real-World Example:

```
user@acme.com → Company "Acme Corp" (auto-created)  
user@techcorp.io → Company "TechCorp" (separate instance)
```

1.5.2 2. AI-Powered Company Profiling

Business Value: 10-minute setup instead of hours of manual data entry.

16 LinkedIn-Style Fields:

Category	Fields
Basic	Name, Tagline, Industry, Founded Year, Size, HQ, Type
About	Description, Mission, Vision, Culture, Values
Business	Products/Services, Target Market, Competitive Advantage, Specialties
Social	LinkedIn, Twitter, Facebook, Logo

AI Process:

1. Enter company website URL
2. AI scrapes multiple pages (home, about, careers)
3. Extracts structured data from JSON-LD and meta tags
4. Generates comprehensive profile in 30 seconds

Technical Stack:

- BeautifulSoup for web scraping
 - OpenAI GPT-4o-mini for intelligent extraction
 - Custom prompts for each field type
 - Re-generation with custom instructions
-

1.5.3 3. Comprehensive Job Descriptions

Business Value: Professional job postings in 2 minutes instead of hours.

11 Detailed Fields:

- Location, Employment Type, Salary Range
- Responsibilities (AI-generated bullets)
- Required Qualifications
- Preferred Qualifications
- Benefits
- Team Information
- Growth Opportunities
- Application Process
- Remote Policy

Smart Features:

- **Context-Aware:** Uses your company's mission and values
- **Regeneration:** Update anytime with new AI content
- **Fine-Tuning:** Add custom instructions ("emphasize remote-first")

Example Workflow:

1. Type: "Senior DevOps Engineer"
 2. AI generates full description using company context
 3. Review & edit if needed
 4. Publish → Auto-match existing candidates
-

1.5.4 4. Enterprise Security & SSO

Authentication Methods:

- Email/Password (Argon2 hashed)
- Microsoft Azure AD (OAuth 2.0)
- Google Workspace (OAuth 2.0)
- Email Verification
- Password Reset (time-limited tokens)

Security Features:

Feature	Implementation
JWT Auth	HS256, 30-min expiry
Password Hashing	Argon2 (memory-hard)
RBAC	5 levels (Super Admin → Interviewer)
Multi-Tenancy	Logical isolation via <code>company_id</code>
Network Isolation	Internal Docker network
CORS	Strict origin policies
SQL Injection	SQLAlchemy ORM
Input Validation	Pydantic schemas
File Security	Unique filenames, path sanitization

Audit Logging:

- Every action logged with user, timestamp, IP
 - Searchable and exportable for compliance
 - Real-time activity feed for admins
-

1.5.5 5. AI-Powered CV Parsing

Business Value: Extract structured data from messy resumes automatically.

Extracted Fields:

- Personal: Name, Email, Phone, Address, Links
- Professional: Summary, Skills, Years of Experience
- Work History: Company, Title, Duration, Description
- Education: School, Degree, Year
- Metadata: Current/Expected Salary

Smart Features:

- **Hidden Link Detection:** Finds LinkedIn/GitHub in PDF metadata
- **Contact Normalization:** Cleans phone numbers and emails
- **Timeline Grouping:** Groups multiple roles at same company
- **Error Recovery:** Handles poorly formatted resumes

Processing Flow:

Upload → Celery Queue → AI Parsing → Database → ChromaDB Embedding

Performance:

- Async processing (doesn't block UI)
 - Batch operations (process 50+ CVs)
 - Retry logic for failed jobs
 - Real-time status updates
-

1.5.6 6. Semantic AI Search

Business Value: Find candidates by meaning, not just keywords.

Features:

- Natural language queries: “Find React developers with AWS experience”
- Auto-sync on startup (self-healing)
- Vector embeddings via OpenAI
- ChromaDB for semantic matching

Technical Architecture:

```
CV Parsed → OpenAI Embedding → ChromaDB Index
      ↓
Search Query → OpenAI Embedding → Vector Search → Ranked Results
```

Self-Healing:

- Detects missing embeddings on startup
 - Automatically re-indexes
 - Manual bulk re-index available
-

1.5.7 7. Candidate Pipeline Management

Business Value: Visual Kanban workflow for managing candidates.

Stages: New → Screening → Interview → Offer → Hired

Features:

- Drag-and-drop candidates between stages
- Bulk actions (assign, delete, reprocess)
- Real-time updates across users
- Version-based cache invalidation
- Mobile-responsive (!!)

Smart Matching:

- As soon as job created → system scans ALL existing CVs
 - Suggests “Silver Medalists” (qualified but not hired)
 - Match score based on skills and experience
-

1.5.8 8. Interview Management

Business Value: Complete interview workflow from scheduling to feedback.

Features:

- Interview assignment to team members
- Multi-stage tracking (Screening, Technical, Final)
- Detailed feedback with ratings (1-10)
- Interview history timeline
- Rescheduling & reassignment
- No-show tracking
- Email notifications

Dedicated Interview Mode:

- Split-screen view (CV left, feedback right)
- Toggle PDF/Parsed view

- Inline scheduling
- Mobile-optimized

Email Notifications:

- Auto-send calendar invites
 - Interview reminders
 - Reschedule notifications
 - Async email delivery (doesn't block UI)
-

1.5.9 9. Full Mobile Responsiveness

Business Value: Recruit on-the-go from any device.

Features:

- Complete feature parity with desktop
- Touch-optimized (44px+ targets)
- Smooth drawer animations
- Hamburger navigation
- Responsive tables and cards
- Mobile-first interview mode

Technical: Tailwind CSS breakpoints for responsive design

1.5.10 10. Real-Time Analytics

Business Value: Data-driven hiring decisions.

Dashboard Metrics:

- Active Jobs
- Total Candidates
- Hires This Month
- Silver Medalists (runner-ups)

Pipeline Insights (per job):

- Average years of experience
- Average current salary
- Average expected salary
- Candidate distribution by stage

Department Stats:

- Headcount by department
- Active jobs per department
- Hiring velocity

1.5.11 11. Background Processing

Business Value: UI remains fast even during heavy operations.

Celery Worker Tasks:

- CV parsing (AI-heavy)
- Bulk candidate reprocessing
- Email sending
- Vector embedding generation

Technology:

- Redis as message broker
 - Automatic retries on failure
 - Task status tracking
 - Real-time progress updates
-

1.5.12 12. AI Department Generator

Business Value: Professional department profiles in seconds.

Features:

- One-click AI generation
- Department description
- Technology stack suggestions
- Job templates for common roles
- Context-aware (uses company profile)

Example:

Department: Engineering

- AI generates:
 - Description: "Building scalable cloud infrastructure..."
 - Tech Stack: React, Python, Kubernetes, AWS
 - Templates: Senior Engineer, DevOps Lead, etc.
-

1.5.13 13. Role-Based Access Control (RBAC)

Business Value: Appropriate access for each team member.

Roles & Permissions:

Role	Scope	Key Permissions
Super Admin	Platform-wide	All companies, all features
Admin	Company-wide	All jobs, billing, user management
Recruiter	Company-wide	Create/manage jobs, move candidates, view-only settings
Hiring Manager	Department only	Manage own department jobs, add interviewers to own dept
Interviewer	Assigned interviews	Submit feedback, salary masked

Department Scoping:

- Hiring Managers only see their department
- Auto-assignment when adding team members
- Cannot access other departments' data

Salary Masking:

- Interviewers see “CONFIDENTIAL” instead of salary
- Prevents conflict of interest

1.5.14 14. Resume Processing Features

Business Value: Fix AI errors and maintain data quality.

Resume/Reprocess Button:

- Visible to: Admin, Recruiter, Hiring Manager
- Triggers: Re-parses ALL pending/failed CVs
- Use case: After uploading bulk CVs, resume interrupted processing

Edit Mode:

- Fix AI parsing errors directly in UI
- Toggle on/off per field
- Update name, experience, summary

Bulk Operations:

- Select 50+ candidates
- Reprocess all selected
- Delete all selected
- Real-time status updates

1.6 Security Architecture

1.6.1 Authentication Flow

```
sequenceDiagram
    User->>Frontend: Login
    Frontend->>Backend: POST /auth/login
    Backend->>Database: Verify credentials
    Database->>Backend: User data
    Backend->>Backend: Generate JWT
    Backend->>Frontend: JWT Token
    Frontend->>Frontend: Store token (localStorage)
    Frontend->>Backend: API calls with token
    Backend->>Backend: Validate JWT
    Backend->>Frontend: Protected data
```

1.6.2 Multi-Tenancy Security

Logical Isolation:

```
# Every query automatically filtered
candidates = db.query(CV).filter(
    CV.company_id == current_user.company_id
).all()
```

Middleware Protection:

- JWT validation on every request
- Company ID extracted from user token
- Automatic filtering applied

Database Design:

- `company_id` foreign key on: CVs, Jobs, Applications, Interviews
- Index on `company_id` for performance
- CASCADE delete ensuring no orphaned data

1.6.3 Network Security

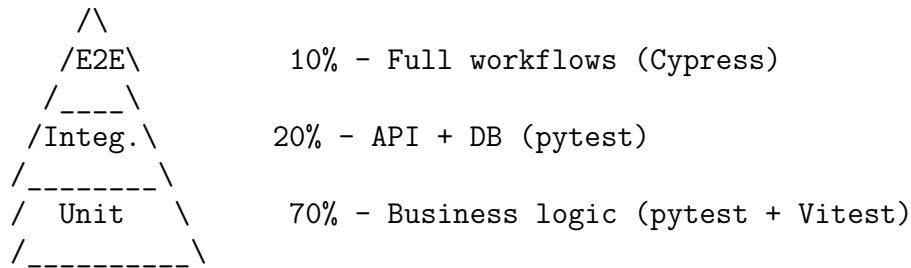
```
Internet
  ↓
  [Nginx Proxy] (optional)
  ↓
  [Docker Network]
    Frontend (exposed: 30004)
    Backend (exposed: 30001)
    [Internal Network]--- Postgres, Redis, ChromaDB
```

Security Layers:

1. Only Frontend & API exposed to host
 2. Database & Redis on internal network
 3. CORS prevents unauthorized origins
 4. Environment variables for secrets
-

1.7 CI/CD & Testing Strategy

1.7.1 Testing Pyramid



1.7.2 Test Coverage (Dec 2025)

Component	Coverage	Status
Backend	76%	Healthy
Frontend	33%	Improving

Backend: 100% coverage for Models, Schemas, Auth

1.7.3 Testing Infrastructure

Backend Unit Tests:

- Tool: pytest
- Database: In-memory SQLite
- Fixtures: `authenticated_client` with pre-configured admin
- Isolation: Database reset per test

Frontend Unit Tests:

- Tool: Vitest + React Testing Library
- Mocks: API responses
- Focus: Component rendering & interactions

E2E Tests:

- Tool: Cypress
- Stack: Full dockerized environment

- Database: Real PostgreSQL with seeded data
 - Zero retries needed (reliable with real services)
-

1.7.4 CI/CD Pipeline

GitHub Actions Workflows:

```
on: [push, pull_request]

jobs:
  test-backend:
    - Setup Python 3.13
    - Install dependencies
    - Lint with Ruff
    - Run pytest (unit tests)

  test-frontend:
    - Setup Node 18
    - Install dependencies
    - Lint with ESLint
    - Run Vitest

  build-and-push:
    needs: [test-backend, test-frontend]
    - Build Docker images
    - Push to Docker Hub
```

1.7.4.1 1. Main CI/CD (cid-cd.yml) Quality Gates:

- Zero test failures
- Zero linting errors
- Successful Docker builds

```
on: [push to main, PR, manual trigger]

steps:
  1. Start E2E stack (docker-compose.e2e.yml)
  2. Wait for services healthy
  3. Reset database schema
  4. Run Alembic migrations
  5. Seed test data
  6. Execute Cypress tests
```

7. Upload artifacts (videos/screenshots)
8. Cleanup

1.7.4.2 2. E2E Tests (`e2e-tests.yml`) E2E Stack Features:

- Isolated environment (separate DB, ports)
 - Real services (no mocks)
 - Automatic cleanup
 - Artifact retention (7 days)
-

1.7.5 Linting & Code Quality

Backend (Python):

- Tool: Ruff (fast Rust-based linter)
- Rules: PEP 8 compliance
- Auto-fix available

Frontend (JavaScript):

- Tool: ESLint
- Rules: React best practices
- Hooks: Pre-commit validation

Command Examples:

```
# Backend
docker compose exec backend ruff check .
docker compose exec backend ruff check . --fix

# Frontend
docker compose exec frontend npm run lint
```

1.7.6 Database Migrations

Tool: Alembic

Workflow:

1. Create migration: `alembic revision -m "description"`
2. Edit generated file (upgrade/downgrade functions)
3. Apply: `alembic upgrade head`
4. Rollback: `alembic downgrade -1`

Best Practices:

- Idempotent migrations (check before alter)

- Descriptive names
- Test both upgrade and downgradeExample migration:

```
def upgrade():
    # Check if column exists before adding
    conn = op.get_bind()
    inspector = sa.inspect(conn)
    columns = [c['name'] for c in inspector.get_columns('interviews')]

    if 'created_at' not in columns:
        op.add_column('interviews',
                      sa.Column('created_at', sa.DateTime(timezone=True)))
```

1.8 Roadmap & Future Features

1.8.1 Current: v1.11.0 (Stable)

Delivered Features:

- Multi-company support
 - AI Company Profiling
 - AI Job Generation
 - Microsoft + Google SSO
 - Email Verification & Password Reset
 - Semantic Search (ChromaDB)
 - Interview Management
 - Mobile Responsive
 - Department Management
 - Role-Based Permissions
-

1.8.2 v2.0 Roadmap (20 weeks)

1.8.2.1 Phase 1: Foundation (Weeks 1-4)

- Google Sign-In (COMPLETE)
- Enhanced Audit Logging
- Department-scoped RBAC improvements

1.8.2.2 Phase 2: Calendar & Automation (Weeks 5-8)

- Google Calendar Integration
- Microsoft Outlook Integration
- Auto Interview Invitations (with attendee rules)

1.8.2.3 Phase 3: Interview Redesign (Weeks 9-12)

- Visual timeline UI
- Structured scorecards
- 360° feedback view
- Pipeline-Interview merge (unified Kanban)

1.8.2.4 Phase 4: Intelligence (Weeks 13-16)

- GraphRAG Knowledge Graph
- Skills Matrix
- Visual network exploration (D3.js)
- Advanced semantic search
- Role-Skill mapping

1.8.2.5 Phase 5: Admin & Analytics (Weeks 17-20)

- Advanced Admin Dashboard
 - Real-time activity feed
 - Time-to-hire analytics
 - Custom report builder
-

1.8.3 Strategic Features Explained

1.8.3.1 Calendar Integration Business Impact: Zero manual scheduling, auto-sync with Google/Outlook

Features:

- Two-way sync
- Availability detection
- Auto-invites to interviewer, hiring manager, recruiter
- Reschedule propagation

1.8.3.2 GraphRAG & Skills Matrix Business Impact: Intelligent candidate discovery and skill gap analysis

Capabilities:

- Knowledge graph: Candidates Skills Projects Companies
- Visual network exploration
- Semantic queries: “Find React developers with cloud experience”
- Skills matrix comparison (candidate vs. job requirements)

Example Use Case:

Query: "Find candidates similar to John Smith"

- Graph algorithm finds:
 - Similar tech stacks
 - Overlapping companies
 - Related project types
-

1.9 Technical Onboarding

1.9.1 Development Setup

Prerequisites:

- Docker & Docker Compose
- OpenAI API Key
- (Optional) SSO credentials

Quick Start:

```
# 1. Clone repository
git clone <repo-url>
cd Headhunter

# 2. Configure .env
cp .env.example .env
# Edit .env with your OPENAI_API_KEY

# 3. Launch stack
docker compose up -d --build

# 4. Run migrations
docker compose exec backend alembic upgrade head

# 5. Access
# Frontend: http://localhost:30004
# API Docs: http://localhost:30001/docs
```

1.9.2 Environment Variables

Required:

```
OPENAI_API_KEY=sk-proj-xxxxx
POSTGRES_USER=user
POSTGRES_PASSWORD=password
DATABASE_URL=postgresql://user:password@db:5432/headhunter_db
```

Optional SSO:

```
# Microsoft
SSO_CLIENT_ID=xxxxx
SSO_CLIENT_SECRET=xxxxx
SSO_TENANT_ID=xxxxx

# Google
GOOGLE_CLIENT_ID=xxxxx
GOOGLE_CLIENT_SECRET=xxxxx
```

Optional Email:

```
MAIL_USERNAME=your-email@example.com
MAIL_PASSWORD=app-password
MAIL_SERVER=smtp.gmail.com
MAIL_PORT=587
```

1.9.3 Key Development Commands

Backend:

```
# Run tests
docker compose exec backend pytest tests/ -v

# Lint
docker compose exec backend ruff check .

# Shell access
docker compose exec backend bash

# View logs
docker compose logs backend -f
```

Frontend:

```
# Run tests
docker compose exec frontend npm run test

# Lint
docker compose exec frontend npm run lint

# Shell access
docker compose exec frontend sh
```

Database:

```

# Create migration
docker compose exec backend alembic revision -m "description"

# Apply migrations
docker compose exec backend alembic upgrade head

# Rollback
docker compose exec backend alembic downgrade -1

# Connect to DB
docker compose exec db psql -U user -d headhunter_db

```

1.9.4 Code Organization

Backend Conventions:

- `api/v1/` - REST endpoints (grouped by entity)
- `models/` - SQLAlchemy models (database tables)
- `services/` - Business logic (AI, search, etc.)
- `core/` - Infrastructure (DB, auth, email)

Frontend Conventions:

- `pages/` - Page-level components
- `components/` - Reusable UI components
- `hooks/` - Custom React hooks
- `services/` - API client functions

Naming:

- Python: `snake_case` (files, functions, variables)
 - JavaScript: `camelCase` (functions, variables), `PascalCase` (components)
 - Database: `snake_case` (tables, columns)
-

1.9.5 Common Tasks

Add New Endpoint:

1. Create Pydantic schemas
2. Add endpoint to `api/v1/`
3. Update API router
4. Write tests
5. Update documentation

Add Database Table:

1. Create SQLAlchemy model in `models/models.py`
2. Generate migration: `alembic revision -m "add_table"`
3. Edit migration file
4. Apply: `alembic upgrade head`
5. Update schemas

Add Frontend Page:

1. Create component in `pages/`
 2. Add route in `App.jsx`
 3. Create API service function
 4. Add navigation link
 5. Style with Tailwind
-

1.10 Business Value Proposition

1.10.1 Cost Comparison

Solution	Annual Cost	Limitations
Greenhouse	\$6,500 - \$18,000	Per-seat pricing, vendor lock-in
Lever	\$12,000 - \$25,000	Cloud-only, limited customization
Workday	\$50,000+	Enterprise only, complex setup
Headhunter AI	\$0 - \$1,200	Open-source, unlimited users

Cost Breakdown (Headhunter AI):

- Software: \$0 (self-hosted)
- OpenAI API: ~\$50-100/month (varies)
- Infrastructure: \$100-1000/month (own servers/cloud)
- **Total:** \$1,800 - \$13,200/year

1.10.2 ROI Calculation

For 50-person company:

Metric	Before	After	Savings
Time per CV	15 min	2 min	87% faster
CVs per month	200	200	-
Hours saved/month	-	43 hrs	~\$2,150/mo @ \$50/hr
Annual savings	-	-	\$25,800/year

Additional Benefits:

- No vendor lock-in
 - Full data ownership
 - GDPR/compliance friendly (self-hosted)
 - Unlimited customization
 - No per-seat fees
-

1.10.3 Strategic Advantages

1. Data Ownership

- All candidate data on your infrastructure
- No third-party data sharing
- Easy GDPR compliance
- Export anytime

2. Customization

- Full source code access
- White-label potential
- Custom integrations
- Tailored workflows

3. Scalability

- Unlimited users
- Unlimited companies (multi-tenant)
- Horizontal scaling (add more workers)
- No licensing constraints

4. Security

- Self-hosted = no cloud exposure
 - Air-gapped deployment possible
 - Custom authentication
 - Audit trails
-

1.10.4 Target Use Cases

Best For:

- Tech-savvy companies with DevOps
- Companies with sensitive data requirements
- High-volume hiring (100+ CVs/month)
- Multi-office/multi-brand organizations
- Cost-conscious startups

Not Ideal For:

- Non-technical teams without IT support
 - Companies preferring SaaS simplicity
 - Very small teams (< 5 hires/year)
-

1.11 Additional Resources

1.11.1 Documentation

- README: /README.md
- Roadmap: /docs/wiki/ROADMAP_V2.md
- Testing Strategy: /docs/wiki/TESTING_STRATEGY.md
- Role Permissions: /docs/wiki/ROLE_PERMISSIONS.md
- Design System: /docs/DESIGN_SYSTEM.md
- Verification: /docs/VERIFICATION.md

1.11.2 API Documentation

- Swagger UI: <http://localhost:30001/docs>
- ReDoc: <http://localhost:30001/redoc>

1.11.3 Support

- GitHub Issues: For bug reports
 - GitHub Discussions: For Q&A
 - Team Wiki: Internal knowledge base
-

1.12 Key Takeaways

1.12.1 Technical Excellence

Modern tech stack (React, FastAPI, Docker)
76% test coverage backend
Automated CI/CD pipeline
Production-ready architecture

1.12.2 Business Value

80% reduction in manual work
\$25K+ annual savings
Complete data ownership
Unlimited scalability

1.12.3 Security

Enterprise-grade authentication
Multi-tenant isolation
RBAC with 5 role levels
SOC 2 ready architecture

1.12.4 Innovation

AI-powered parsing & matching
Semantic search capabilities
Context-aware job generation
Mobile-first design

1.13 Next Steps

1.13.1 For Development Team

1. Complete Phase 1 roadmap (RBAC improvements)
2. Begin calendar integration POC
3. Increase frontend test coverage to 60%

1.13.2 For Business Team

1. Define pricing model (if productizing)
2. Create customer onboarding materials
3. Develop case studies

1.13.3 For Operations

1. Setup production infrastructure
 2. Configure monitoring (Sentry, DataDogConfigure backup strategy)
 3. Implement disaster recovery
-

Built with by the Headhunter AI Engineering Team

This presentation document is markdown-based and can be easily converted to PDF using Pandoc or to PowerPoint using Marp or similar tools.

1.14 Conversion Instructions

1.14.1 To PDF

```
# Using Pandoc  
pandoc presentation.md -o presentation.pdf --pdf-engine=xelatex  
  
# Using Marp  
marp presentation.md -o presentation.pdf
```

1.14.2 To PowerPoint

```
# Using Pandoc  
pandoc presentation.md -o presentation.pptx  
  
# Using Marp  
marp presentation.md -o presentation.pptx
```

1.14.3 To HTML Slides

```
# Using Marp  
marp presentation.md -o presentation.html --theme default  
  
# Using reveal.js via Pandoc  
pandoc presentation.md -t revealjs -s -o presentation.html
```